Subtask 10.2.3

How do you think the DBMS would execute this query?

```
EXPLAIN EXTENDED
SELECT COUNT(*) FROM Orders;
```

#	id	select_type	table	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	1	SIMPLE	Orders	index	NULL	CustomerID	5	NULL	861	100.00	Using index

Only the index gets checked instead of the whole tuple.

Is this different from what you expected?

Yes, it is very different to what I expected. I thought it would read through each tuple.

Can you reason why this strategy was chosen?

It is faster to check just the index rather than checking each tuple fully.